

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sb
```

```
sat=pd.read_csv('https://github.com/ybifoundation/Dataset/raw/main/SAT%20GPA.csv')
```

```
sat.head()
```

	SAT	GPA
0	1270	3.4
1	1220	4.0
2	1160	3.8
3	950	3.8
4	1070	4.0

```
sat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    SAT      1000 non-null    int64
1    GPA      1000 non-null    float64
dtypes: float64(1), int64(1)
memory usage: 15.8 KB
```

```
sat.describe()
```


	SAT	GPA
count	1000.000000	1000.000000
mean	1033.290000	3.203700
std	142.873681	0.542541
min	530.000000	1.800000

```
sat.corr()
```

	SAT	GPA
SAT	1.000000	0.429649
GPA	0.429649	1.000000

```
from seaborn.axisgrid import pairplot
sb.pairplot(sat)
```

```
<seaborn.axisgrid.PairGrid at 0x7f7d268badd0>
```



```
sat.columns
```

```
Index(['SAT', 'GPA'], dtype='object')
```

```
y=sat['SAT']
```

```
y.shape
```

```
(1000,)
```

```
x=sat[['GPA']]
```

```
x.shape
```

```
(1000, 1)
```

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test=train_test_split(x,y,train_size=0.7,random_state=2529)
```

```
x_train.shape, x_test.shape, y_train.shape, y_test.shape
```

```
((700, 1), (300, 1), (700,), (300,))
```

```
x_train
```

	GPA
669	3.7
583	3.7
688	2.8
422	3.9
825	4.0
...	...
740	2.5
399	2.6
828	3.2
562	2.7
352	3.0

700 rows × 1 columns

```
from sklearn.linear_model import LinearRegression
```

```
reg = LinearRegression()
```

```
reg.fit(x_train,y_train)
```

```
LinearRegression()
```

```
reg.intercept_
```

```
673.2291896122774
```

```
reg.coef_
```

```
array([111.01584994])
```

```
y_pred=reg.predict(x_test)
```

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.1046927663671282
```

```
mean_absolute_error(y_test,y_pred)
```


```
105.93877473699905
```

```
r2_score(y_test,y_pred)
```

```
0.18785383761597474
```

```
sb.regplot(x='GPA',y='SAT',data=sat)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7d1f5c5610>
```



```
import matplotlib.pyplot as plt
```

```
df=pd.read_csv('https://github.com/ybifoundation/Dataset/raw/main/Boston.csv')
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	23.1
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	17.8
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	17.8
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	18.7
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	18.7

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM        506 non-null    float64
1    ZN          506 non-null    float64
2    INDUS       506 non-null    float64
3    CHAS        506 non-null    int64
4    NX          506 non-null    float64
5    RM          506 non-null    float64
6    AGE         506 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    int64
9    TAX         506 non-null    float64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       506 non-null    float64
13   MEDV       506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
df.describe()
```

	CRIM	ZN	INDUS	CHAS	NX	RM	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.00
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.57
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.14
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.90

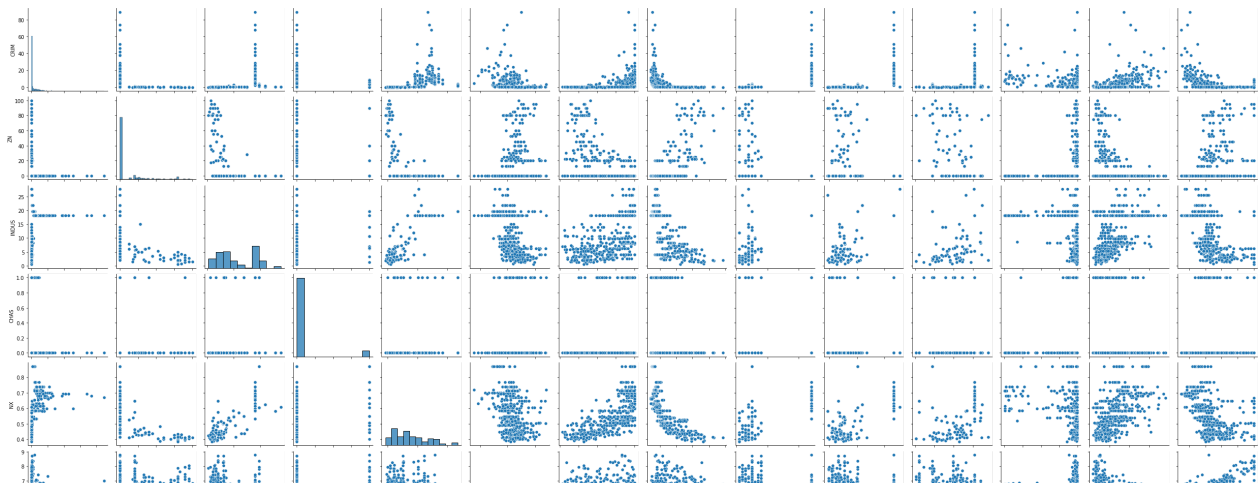
```

import seaborn as so

50%    0.256510    0.000000    9.690000    0.000000    0.538000    6.208500    77.50

so.pairplot(df)

```



```
df.columns
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
      'PTRATIO', 'B', 'LSTAT', 'MEDV'],  
      dtype='object')
```



```
y=df['MEDV']
```



```
x=df[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'B']]
```



```
x.shape
```

```
(506, 11)
```



```
from sklearn.preprocessing import StandardScaler
```



```
sc=StandardScaler()
```

```
x_train=sc.fit_transform(x_train)
```

```
x_test=sc.fit_transform(x_test)
```

```
x_train
```

```
array([[ 0.9146784 ],  
       [ 0.9146784 ],  
       [-0.75077649],  
       [ 1.28477948],  
       [ 1.46983002],  
       [ 0.17447623],  
       [ 1.09972894],  
       [-0.3806754 ],  
       [ 1.09972894],  
       [-2.23118083],
```

```
[-2.04613029],  
[-2.04613029],  
[ 0.17447623],  
[ 0.9146784 ],  
[-2.60128191],  
[-0.3806754 ],  
[-0.75077649],  
[ 0.72962785],  
[ 1.46983002],  
[ 0.9146784 ],  
[ 1.28477948],  
[ 0.17447623],  
[ 1.28477948],  
[-0.01057432],  
[ 0.9146784 ],  
[-0.56572594],  
[ 1.09972894],  
[-1.6760292 ],  
[ 0.17447623],  
[ 1.09972894],  
[-0.3806754 ],  
[-0.3806754 ],  
[ 1.09972894],  
[ 1.09972894],  
[-0.93582703],  
[ 0.54457731],  
[-0.01057432],  
[ 0.35952677],  
[ 0.17447623],  
[ 0.9146784 ],  
[ 0.54457731],  
[ 1.09972894],  
[-1.30592812],  
[ 0.9146784 ],  
[-1.6760292 ],  
[-0.93582703],  
[-1.30592812],  
[ 1.46983002],  
[-0.19562486],  
[ 0.54457731],  
[-0.3806754 ],  
[ 1.46983002],  
[ 1.46983002],  
[ 1.09972894],  
[-0.3806754 ],  
[ 0.54457731],  
[-1.86107974],  
[ 0.35952677],
```

```
from sklearn.linear_model import LinearRegression
```

```
model.fit(x_train,y_train)
```

```
LinearRegression()
```



```
model.intercept_
```

```
1029.1142857142856
```

```
model.coef_
```

```
array([59.99217745])
```

```
model=LinearRegression()
```

```
df.columns
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
      'PTRATIO', 'B', 'LSTAT', 'MEDV'],  
      dtype='object')
```

```
y_pred=model.predict(x_test)
```

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.1046927663671282
```

```
mean_absolute_error(y_test,y_pred)
```

```
105.8772997993962
```

```
r2_score(y_test,y_pred)
```

```
0.18858714592440917
```

```
df1=pd.read_csv('https://github.com/ybifoundation/Dataset/raw/main/Fish.csv')
```

```
df1.head()
```

	Category	Species	Weight	Height	Width	Length1	Length2	Length3
0	1	Brown	340.0	44.5000	4.0000	22.0	25.4	22.0

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 159 entries, 0 to 158
Data columns (total 8 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Category    159 non-null    int64
1   Species     159 non-null    object
2   Weight      159 non-null    float64
3   Height      159 non-null    float64
4   Width       159 non-null    float64
5   Length1     159 non-null    float64
6   Length2     159 non-null    float64
7   Length3     159 non-null    float64
dtypes: float64(6), int64(1), object(1)
memory usage: 10.1+ KB
```

```
df1.describe()
```

	Category	Weight	Height	Width	Length1	Length2	Length3
count	159.000000	159.000000	159.000000	159.000000	159.000000	159.000000	159.000000
mean	3.264151	398.326415	8.970994	4.417486	26.247170	28.415723	31.247170
std	1.704249	357.978317	4.286208	1.685804	9.996441	10.716328	11.632817
min	1.000000	0.000000	1.728400	1.047600	7.500000	8.400000	8.800000
25%	2.000000	120.000000	5.944800	3.385650	19.050000	21.000000	23.100000
50%	3.000000	273.000000	7.786000	4.248500	25.200000	27.300000	29.400000
75%	4.500000	650.000000	12.365900	5.584500	32.700000	35.500000	39.600000
max	7.000000	1650.000000	18.957000	8.142000	59.000000	63.400000	68.000000

```
df1.columns
```

```
Index(['Category', 'Species', 'Weight', 'Height', 'Width', 'Length1',
      'Length2', 'Length3'],
      dtype='object')
```

```
df1.shape
```

```
(159, 8)
```

```
y=df1['Weight']
```

```
y.shape
```

```
(159,)
```

```
y
```

```
0      242.0
```

```
1      290.0
```

```
2      340.0
```

```
3      363.0
```

```
4      430.0
```

```
...
```

```
154     12.2
```

```
155     13.4
```

```
156     12.2
```

```
157     19.7
```

```
158     19.9
```

```
Name: Weight, Length: 159, dtype: float64
```

```
x=[['Weight', 'Height', 'Width', 'Length1',  
    'Length2', 'Length3']]
```

```
x=df1.drop(['Species', 'Weight'], axis=1)
```

```
x.shape
```

```
(159, 6)
```

```
x
```

	Category	Height	Width	Length1	Length2	Length3
0	1	11.5200	4.0200	23.2	25.4	30.0
1	1	12.4800	4.2050	24.0	26.2	31.0

```
df_new=df1.sample(1)
```

```
df_new
```

	Category	Species	Weight	Height	Width	Length1	Length2	Length3
41	5	Roach	110.0	6.1677	3.3957	19.1	20.8	23.1
157	6	2.8728	2.0672	13.2	14.3	15.2		

```
x_new=df_new[['Weight', 'Height', 'Length1',  
              'Length2', 'Length3']]
```

```
x_new.shape
```

```
(1, 5)
```

```
from sklearn.linear_model import LinearRegression
```

```
model1=LinearRegression()
```

```
model1.fit(x_test,y_test)
```

```
LinearRegression()
```

```
x_new
```

	Weight	Height	Length1	Length2	Length3
41	110.0	6.1677	19.1	20.8	23.1

```
model1.coef_
```

```
array([64.62382312])
```

```
model1.intercept_
```

```
1043.0333333333333
```


```
import statsmodels.api as sm
```

```
/usr/local/lib/python3.7/dist-packages/statsmodels/tools/_testing.py:19: FutureWarning:
import pandas.util.testing as tm
```



```
x=sm.add_constant(x)
```

```
/usr/local/lib/python3.7/dist-packages/statsmodels/tsa/tsatools.py:117: FutureWarning:
x = pd.concat(x[:,order], 1)
```



```
x.head()
```

	const	Category	Height	Width	Length1	Length2	Length3
0	1.0	1	11.5200	4.0200	23.2	25.4	30.0
1	1.0	1	12.4800	4.3056	24.0	26.3	31.2
2	1.0	1	12.3778	4.6961	23.9	26.5	31.1
3	1.0	1	12.7300	4.4555	26.3	29.0	33.5
4	1.0	1	12.4440	5.1340	26.5	29.0	34.0

```
sample=sm.OLS(y_test,x_test).fit()
```

```
y_pred=sample.predict(y_train)
```

```
y_pred
```

```
669    55192.803252
583    63591.708095
688    62391.864546
422    70190.847615
825    76789.987134
...
740    55192.803252
399    53992.959704
828    55192.803252
562    50393.429057
352    64791.551644
Length: 700, dtype: float64
```

```
y_pred.shape
```

```
(700,)
```

```
from sklearn.metrics import mean_squared_error
```

```
mean_squared_error(y_test,y_pred)
```

```
1104767.546960686
```

```
print(sample.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  SAT    R-squared (uncentered):          0.003
Model:                          OLS    Adj. R-squared (uncentered):       0.002
Method:                        Least Squares    F-statistic:                2.339
Date:                          Tue, 19 Apr 2022    Prob (F-statistic):          0.127
Time:                          15:46:45    Log-Likelihood:              -5854.2
No. Observations:                700    AIC:                        1.171e+04
Df Residuals:                    699    BIC:                        1.171e+04
Df Model:                        1
Covariance Type:                nonrobust
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
x1              59.9922      39.226        1.529      0.127      -17.023      137.008
=====
Omnibus:                        0.718    Durbin-Watson:                  0.032
Prob(Omnibus):                  0.698    Jarque-Bera (JB):                0.796
Skew:                          0.071    Prob(JB):                        0.672
Kurtosis:                      2.914    Cond. No.                        1.00
=====

```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specif



```
sample
```

```
<statsmodels.regression.linear_model.RegressionResultsWrapper at 0x7f7d1d5e12d0>
```

