

## predictive analytics

```
import pandas as pd
```

```
import numpy as num
```

```
import matplotlib.pyplot as mp
```

```
import seaborn as sn
```

```
df=pd.read_csv("https://github.com/YBI-foundation/Dataset/raw/main/MPG.csv")
```

```
df.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
0	18.0	8	307.0	130.0	3504	12.0	70	usa
1	15.0	8	350.0	165.0	3693	11.5	70	usa

```
df.tail()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin
393	27.0	4	140.0	86.0	2790	15.6	82	usa
394	44.0	4	97.0	52.0	2130	24.6	82	europa
395	32.0	4	135.0	84.0	2295	11.6	82	usa

```
df.nunique()
```

```
mpg          129
cylinders      5
displacement  82
horsepower    93
weight       351
```

```

acceleration    95
model_year      13
origin           3
name            305
dtype: int64

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   mpg              398 non-null   float64
1   cylinders        398 non-null   int64   
2   displacement     398 non-null   float64
3   horsepower       392 non-null   float64
4   weight           398 non-null   int64   
5   acceleration     398 non-null   float64
6   model_year       398 non-null   int64   
7   origin           398 non-null   object  
8   name             398 non-null   object  
dtypes: float64(4), int64(3), object(2)
memory usage: 28.1+ KB

```

```
df.describe()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model
<b>count</b>	398.000000	398.000000	398.000000	392.000000	398.000000	398.000000	398.000000
<b>mean</b>	23.514573	5.454774	193.425879	104.469388	2970.424623	15.568090	76.000000
<b>std</b>	7.815984	1.701004	104.269838	38.491160	846.841774	2.757689	3.600000
<b>min</b>	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70.000000
<b>25%</b>	17.500000	4.000000	104.250000	75.000000	2223.750000	13.825000	73.000000
<b>50%</b>	23.000000	4.000000	148.500000	93.500000	2803.500000	15.500000	76.000000
<b>75%</b>	29.000000	8.000000	262.000000	126.000000	3608.000000	17.175000	79.000000
<b>max</b>	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	82.000000

```
df.corr()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	mo
<b>mpg</b>	1.000000	-0.775396	-0.804203	-0.778427	-0.831741	0.420289	
<b>cylinders</b>	-0.775396	1.000000	0.950721	0.842983	0.896017	-0.505419	.
<b>displacement</b>	-0.804203	0.950721	1.000000	0.897257	0.932824	-0.543684	.
<b>horsepower</b>	-0.778427	0.842983	0.897257	1.000000	0.861538	-0.680106	.

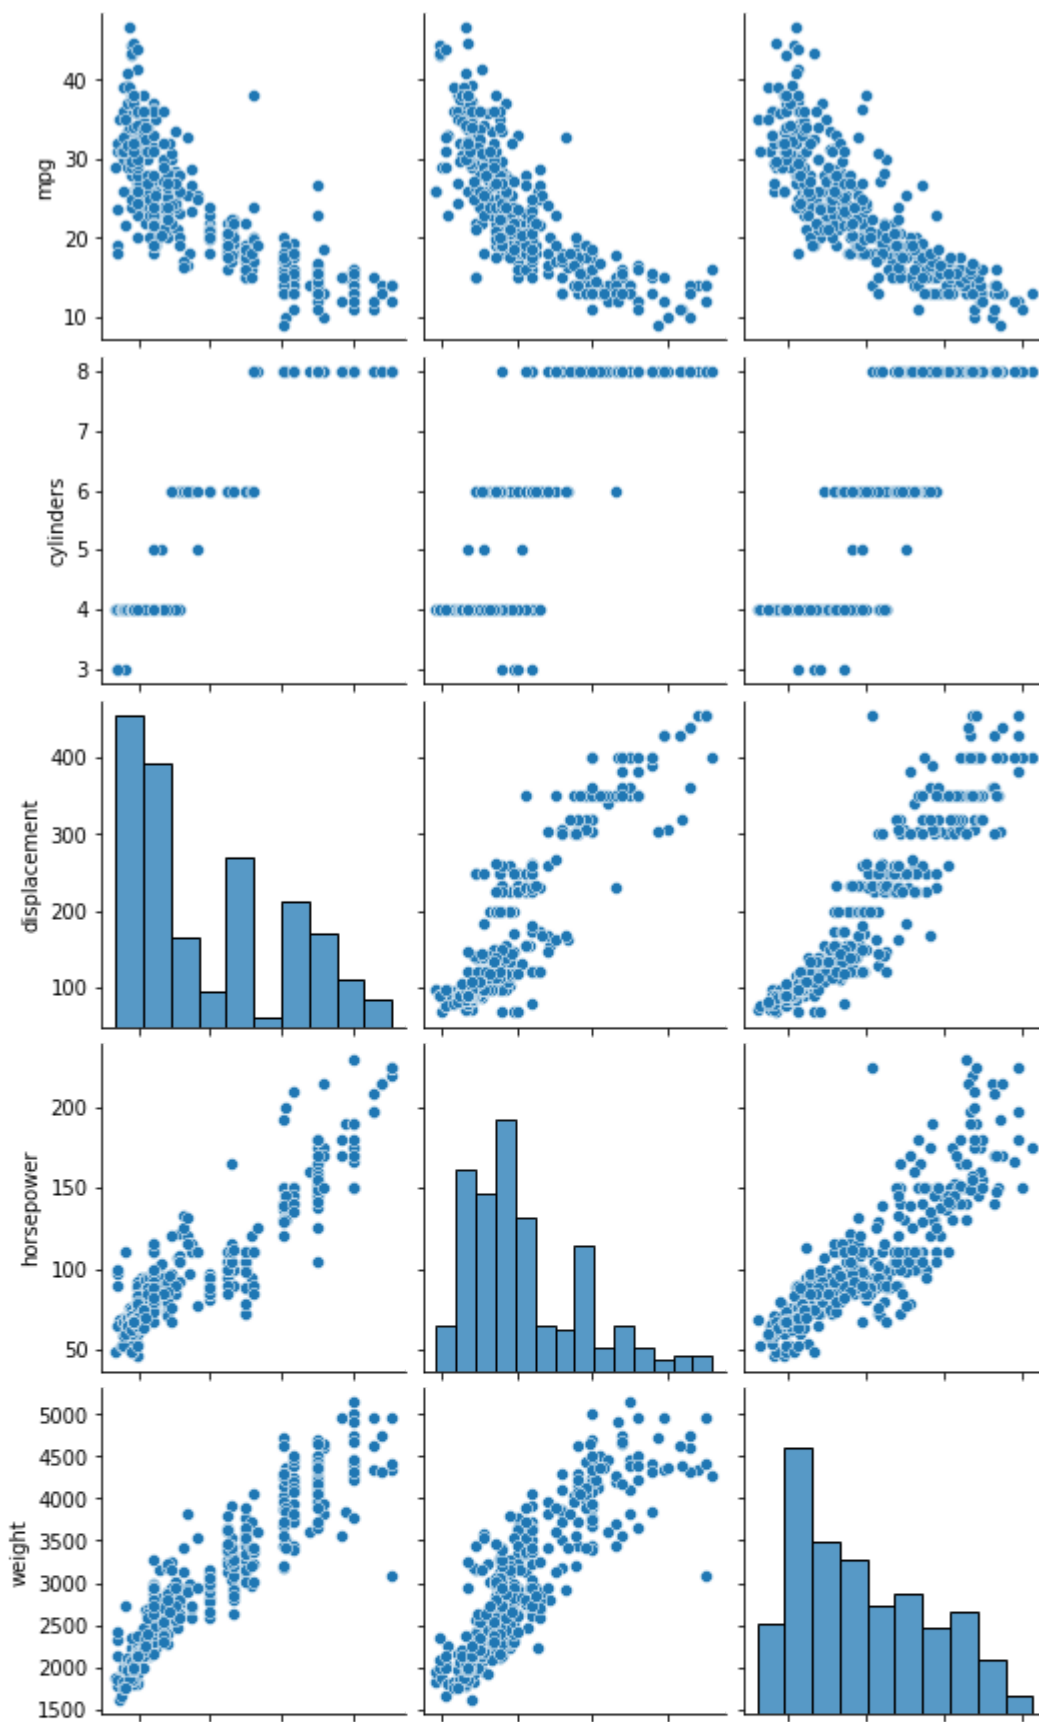
```
df=df.dropna()
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 392 entries, 0 to 397
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   mpg             392 non-null   float64
1   cylinders       392 non-null   int64
2   displacement    392 non-null   float64
3   horsepower      392 non-null   float64
4   weight          392 non-null   int64
5   acceleration    392 non-null   float64
6   model_year      392 non-null   int64
7   origin          392 non-null   object
8   name            392 non-null   object
dtypes: float64(4), int64(3), object(2)
memory usage: 30.6+ KB
```


```
sn.pairplot(df,x_vars=['displacement','horsepower','weight'])
```

```
<seaborn.axisgrid.PairGrid at 0x7f0c93f09ed0>
```



```
sn.regplot(x='displacement',y='mpg',data=df)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f0c8eef9e10>



df.columns

```
Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
      'acceleration', 'model_year', 'origin', 'name'],
      dtype='object')
```

y=df['mpg']

y.shape

(392,)

x=df[['displacement', 'weight', 'acceleration']]

x.shape

(392, 3)

```
from sklearn.preprocessing import StandardScaler
```

```
ss=StandardScaler()
```

```
x=ss.fit_transform(x)
```

x

```
array([[ 1.07728956,  0.62054034, -1.285258  ],
       [ 1.48873169,  0.84333403, -1.46672362],
       [ 1.1825422 ,  0.54038176, -1.64818924],
       ...,
       [-0.56847897, -0.80463202, -1.4304305 ],
       [-0.7120053 , -0.41562716,  1.11008813],
       [-0.72157372, -0.30364091,  1.40043312]])
```

```
pd.DataFrame(x).describe()
```

	0	1	2
<b>count</b>	3.920000e+02	3.920000e+02	3.920000e+02
<b>mean</b>	-2.537653e-16	5.607759e-17	6.117555e-16
<b>std</b>	1.001278e+00	1.001278e+00	1.001278e+00
<b>min</b>	-1.209563e+00	-1.608575e+00	-2.736983e+00
<b>25%</b>	-8.555316e-01	-8.868535e-01	-6.410551e-01
<b>50%</b>	-4.153842e-01	-2.052109e-01	-1.499869e-02
<b>75%</b>	7.782764e-01	7.510927e-01	5.384714e-01
<b>max</b>	2.493416e+00	2.549061e+00	3.360262e+00

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.7,random_state=2535)
```

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
((274, 3), (118, 3), (274,), (118,))
```

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression()
```

```
lr.intercept_
```

```
23.481757374918818
```

```
lr.coef_
```

```
array([-0.3969492 , -5.81652961,  0.84382013])
```

```
y_pred=lr.predict(x_test)
```

y\_pred

```
array([27.0950622 , 24.86846697, 25.46608867, 20.18920635, 31.58350197,
       31.07258189, 25.8960753 , 20.05161656, 14.19571234, 28.80613919,
       25.41127903, 11.52775544, 28.85413497, 20.1402067 , 19.92816833,
       15.46828926, 15.18457187, 18.64192416, 28.15641361, 14.88968591,
       14.47829349, 20.56365413, 20.64130953, 23.37278745, 32.21165547,
       18.17166849, 13.7990552 , 11.86580018, 30.78406874, 32.56912928,
       24.00456213, 29.42837442, 25.558358 , 13.21581137, 21.1979567 ,
       26.89621049, 31.36274567, 28.55664204, 29.46346568, 30.81338038,
       26.16361613, 29.24905037, 23.45861862, 30.65157817, 33.26899217,
       11.93439594, 18.8961787 , 27.11860968, 13.9210924 , 21.54222525,
       11.94238978, 26.00051733, 27.12174109, 16.78845349, 12.60957108,
       27.76932967, 29.71535153, 27.71270789, 23.26653444, 18.82533404,
       14.21400151, 21.57102879, 21.06774658, 15.00213014, 31.22018153,
       23.85502572, 17.20993611, 26.27905746, 15.47922728, 25.31382302,
       24.44682056, 9.49115678, 18.29693791, 30.14639531, 10.58013426,
       23.24593446, 24.84436754, 29.40520557, 26.75682528, 8.54723467,
       32.78899625, 14.62408239, 15.45600422, 29.3864358 , 23.71451515,
       25.06107848, 21.01760441, 29.1670417 , 21.03359794, 29.34433562,
       30.45713615, 30.60236054, 19.30241702, 6.7896799 , 15.94287186,
       12.39166433, 17.10206798, 26.27292266, 24.1721167 , 12.22118426,
       15.69647721, 17.10525464, 31.64640977, 20.11024545, 26.62598022,
       34.05599176, 12.75776101, 30.64782582, 24.3743673 , 26.02748605,
       25.05729718, 28.32368724, 24.46624597, 26.67891839, 31.32791246,
       27.32951181, 29.34477957, 30.31335789])
```

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score
```

```
mean_absolute_error(y_test,y_pred)
```

```
3.2445387480016987
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.14972398521815397
```

```
r2_score(y_test,y_pred)
```

```
0.7088621831979525
```

```
from sklearn.preprocessing import PolynomialFeatures
```

```
poly=PolynomialFeatures(degree=2,interaction_only=True,include_bias=False,order='c')
```

```
x_train2=poly.fit_transform(x_train)
```

```
x_test2=poly.fit_transform(x_test)
```

```
lr.fit(x_train2,y_train)

LinearRegression()

lr.intercept_

21.76889066826468

lr.coef_

array([-1.9108918 , -5.28015114,  1.01166591,  1.67884703, -0.80744566,
        0.80323871])

y_pred_poly=lr.predict(x_test2)

r2_score(y_test,y_pred_poly)

0.7490061639178867

from sklearn.datasets import load_digits

df=load_digits()

df.images.shape

(1797, 8, 8)

df.images[0].shape

(8, 8)

df.images[0]

array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.],
       [ 0.,  0., 13., 15., 10., 15.,  5.,  0.],
       [ 0.,  3., 15.,  2.,  0., 11.,  8.,  0.],
       [ 0.,  4., 12.,  0.,  0.,  8.,  8.,  0.],
       [ 0.,  5.,  8.,  0.,  0.,  9.,  8.,  0.],
       [ 0.,  4., 11.,  0.,  1., 12.,  7.,  0.],
       [ 0.,  2., 14.,  5., 10., 12.,  0.,  0.],
       [ 0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])

n_samples=len(df.images)
data=df.images.reshape((n_samples, -1))
```



```
data[0]
```

```
array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
        15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
        12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
         0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
        10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]
```

```
data.shape
```

```
(1797, 64)
```

```
data[0].shape
```

```
(64,)
```

```
data.min()
```

```
0.0
```

```
data.max()
```

```
16.0
```

```
data[0]
```

```
array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
        15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
        12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
         0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
        10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.]
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf=RandomForestClassifier()
```

```
from sklearn.metrics import confusion_matrix,classification_report
```

```
len(df.images)
```

```
1797
```

car price prediction

```
df1=pd.read_csv("https://github.com/YBI-foundation/Dataset/raw/main/Car%20Price.csv")
```

```
df1.head()
```

	Brand	Model	Year	Selling_Price	KM_Driven	Fuel	Seller_Type	Transmission
0	Maruti	Maruti 800 AC	2007	60000	70000	Petrol	Individual	Manual
1	Maruti	Maruti Wagon R LXI Minor	2007	135000	50000	Petrol	Individual	Manual
Hyundai								

```
df1.tail()
```

	Brand	Model	Year	Selling_Price	KM_Driven	Fuel	Seller_Type	Transmission
4335	Hyundai	Hyundai i20 Magna 1.4 CRDi (Diesel)	2014	409999	80000	Diesel	Individual	Manual
Hyundai i20								

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4340 entries, 0 to 4339
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Brand           4340 non-null   object
1   Model           4340 non-null   object
2   Year            4340 non-null   int64
3   Selling_Price   4340 non-null   int64
4   KM_Driven       4340 non-null   int64
5   Fuel            4340 non-null   object
6   Seller_Type     4340 non-null   object
7   Transmission    4340 non-null   object
8   Owner           4340 non-null   object
dtypes: int64(3), object(6)
memory usage: 305.3+ KB
```

```
df1.describe()
```

	Year	Selling_Price	KM_Driven
<b>count</b>	4340.000000	4.340000e+03	4340.000000
<b>mean</b>	2013.090783	5.041273e+05	66215.777419
<b>std</b>	4.215344	5.785487e+05	46644.102194
<b>min</b>	1992.000000	2.000000e+04	1.000000
<b>25%</b>	2011.000000	2.087498e+05	35000.000000
<b>50%</b>	2014.000000	3.500000e+05	60000.000000
<b>75%</b>	2016.000000	6.000000e+05	90000.000000

```
df1[['Brand']].value_counts()
```

```
Brand
Maruti      1280
Hyundai     821
Mahindra    365
Tata        361
Honda       252
Ford        238
Toyota      206
Chevrolet   188
Renault     146
Volkswagen  107
Skoda       68
Nissan       64
Audi        60
BMW         39
Fiat        37
Datsun      37
Mercedes-Benz 35
Mitsubishi  6
Jaguar      6
Land        5
Ambassador  4
Volvo       4
Jeep        3
OpelCorsa   2
MG          2
Isuzu       1
Force       1
Daewoo      1
Kia         1
dtype: int64
```

```
df1[['Fuel']].value_counts()
```

```
Fuel
Diesel    2153
Petrol    2123
```

```

CNG          40
LPG          23
Electric     1
dtype: int64

```

```
df1[['Model']].value_counts()
```

```

Model
Maruti Swift Dzire VDI          69
Maruti Alto 800 LXI            59
Maruti Alto LXi                47
Hyundai EON Era Plus           35
Maruti Alto LX                 35
..
Mahindra KUV 100 G80 K4 Plus    1
Mahindra KUV 100 mFALCON D75 K8 1
Mahindra KUV 100 mFALCON D75 K8 AW 1
Mahindra KUV 100 mFALCON G80 K2 Plus 1
Volvo XC60 D5 Inscription       1
Length: 1491, dtype: int64

```

```
df1.columns
```

```

Index(['Brand', 'Model', 'Year', 'Selling_Price', 'KM_Driven', 'Fuel',
       'Seller_Type', 'Transmission', 'Owner'],
      dtype='object')

```

```
df1.shape
```

```
(4340, 9)
```

```
df1.replace({'Transmission':{'Manual':0,'Automatic':1}},inplace=True)
```

```
y=df1['Selling_Price']
```

```
y.shape
```

```
(4340,)
```

```
x=df1[['Year','Transmission']]
```

```
x.shape
```

```
(4340, 2)
```

```
x
```

	Year	Transmission
<b>0</b>	2007	0
<b>1</b>	2007	0
<b>2</b>	2012	0
<b>3</b>	2017	0
<b>4</b>	2014	0
...	...	...
<b>4335</b>	2014	0
<b>4336</b>	2014	0
<b>4337</b>	2009	0
<b>4338</b>	2016	0
<b>4339</b>	2016	0

4340 rows × 2 columns

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.5,random_state=2539)
```

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
((2170, 2), (2170, 2), (2170,), (2170,))
```

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression()
```

y

```
0      60000
1     135000
2     600000
3     250000
4     450000
...
4335   409999
4336   409999
```

```
4337    110000
4338    865000
4339    225000
Name: Selling_Price, Length: 4340, dtype: int64
```

```
y_pred=lr.predict(x_test)
```

```
y_pred.shape
```

```
(2170,)
```

```
y_pred
```

```
array([504237.05512911, 357265.74335442, 259284.86883798, ...,
       357265.74335442, 308275.3060962 , 455246.61787088])
```

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error,r2_score,mean_
```

```
mean_absolute_error(y_test,y_pred)
```

```
240430.41301709128
```

```
r2_score(y_test,y_pred)
```

```
0.36870936179375813
```

```
mean_squared_error(y_test,y_pred)
```

```
215527527203.4092
```


```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.6914047191678929
```

```
mp.scatter(y_test,y_pred)
```

```
mp.title('Actual prices vs Predicted prices')
```

```
mp.show()
```



```
df_new=df1.sample(1)
```



```
df_new
```

	Brand	Model	Year	Selling_Price	KM_Driven	Fuel	Seller_Type	Transmission	C
	Tata								

```
df_new.shape
```

(1, 9)

bike price prediction

```
df2=pd.read_csv("https://github.com/YBI-foundation/Dataset/raw/main/Bike%20Prices.csv")
```

```
df2.head()
```

	Brand	Model	Selling_Price	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Pric
0	TVS	TVS XL 100	30000	2017	Individual	1st owner	8000	30490.0
1	Bajaj	Bajaj ct 100	18000	2017	Individual	1st owner	35000	32000.0
2	Yo	Yo Style	20000	2011	Individual	1st owner	10000	37675.0

```
df2.tail()
```

Brand	Model	Selling_Price	Year	Seller_Type	Owner	KM_Driven	Ex_Showroom_Pr
-------	-------	---------------	------	-------------	-------	-----------	----------------

```
df2.describe()
```

	Selling_Price	Year	KM_Driven	Ex_Showroom_Price
<b>count</b>	1061.000000	1061.000000	1061.000000	6.260000e+02
<b>mean</b>	59638.151744	2013.867107	34359.833176	8.795871e+04
<b>std</b>	56304.291973	4.301191	51623.152702	7.749659e+04
<b>min</b>	5000.000000	1988.000000	350.000000	3.049000e+04
<b>25%</b>	28000.000000	2011.000000	13500.000000	5.485200e+04
<b>50%</b>	45000.000000	2015.000000	25000.000000	7.275250e+04
<b>75%</b>	70000.000000	2017.000000	43000.000000	8.703150e+04
<b>max</b>	760000.000000	2020.000000	880000.000000	1.278000e+06

```
df2.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1061 entries, 0 to 1060
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Brand                  1061 non-null  object
1   Model                  1061 non-null  object
2   Selling_Price          1061 non-null  int64
3   Year                   1061 non-null  int64
4   Seller_Type            1061 non-null  object
5   Owner                  1061 non-null  object
6   KM_Driven              1061 non-null  int64
7   Ex_Showroom_Price      626 non-null   float64
dtypes: float64(1), int64(3), object(4)
memory usage: 66.4+ KB
```

```
df2=df2.dropna()
```

```
df2[['Brand']].value_counts()
```

Brand	
Honda	170
Bajaj	143
Hero	108
Yamaha	94
Royal	40
TVS	23
Suzuki	18
KTM	6



```

Mahindra      6
Kawasaki      4
UM            3
Activa        3
Harley        2
Vespa         2
BMW           1
Hyosung       1
Benelli       1
Yo            1
dtype: int64

```

```
df2[['Owner']].value_counts()
```

```

Owner
1st owner    556
2nd owner     66
3rd owner      3
4th owner      1
dtype: int64

```

```
df2['Model'].value_counts()
```

```

Honda Activa [2000-2015]    23
Honda CB Hornet 160R       22
Bajaj Pulsar 180           20
Bajaj Discover 125         16
Yamaha FZ S V 2.0          16
..
TVS Radeon                 1
Hero Ignitor Disc          1
Bajaj V12                  1
Yamaha Cygnus Ray ZR       1
Harley-Davidson Street Bob 1
Name: Model, Length: 183, dtype: int64

```

```
df2.columns
```

```

Index(['Brand', 'Model', 'Selling_Price', 'Year', 'Seller_Type', 'Owner',
      'KM_Driven', 'Ex_Showroom_Price'],
      dtype='object')

```

```
df2.shape
```

```
(626, 8)
```

```
df2.replace({'Seller_Type':{'Individual':0,'Dealer':1}},inplace=True)
```

```
y=df2["Selling_Price"]
```

y

```

0      30000
1      18000
2      20000
3      25000
4      24999
...
621    330000
622    300000
623    425000
624    760000
625    750000
Name: Selling_Price, Length: 626, dtype: int64

```

```
x=df2[["Year","Seller_Type"]]
```

x.shape

```
(626, 2)
```

x

	Year	Seller_Type
<b>0</b>	2017	0
<b>1</b>	2017	0
<b>2</b>	2011	0
<b>3</b>	2010	0
<b>4</b>	2012	0
...	...	...
<b>621</b>	2014	0
<b>622</b>	2011	0
<b>623</b>	2017	0
<b>624</b>	2019	0
<b>625</b>	2013	0

626 rows × 2 columns

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.6,random_state=2532)
```

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
((375, 2), (251, 2), (375,), (251,))
```

```
from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression()
```

```
y
```

```
0      30000
1      18000
2      20000
3      25000
4      24999
...
621    330000
622    300000
623    425000
624    760000
625    750000
Name: Selling_Price, Length: 626, dtype: int64
```

```
y_pred=lr.predict(x_test)
```

```
y.shape
```

```
(626,)
```

```
y_pred
```

```
array([ 62064.35694737,  47153.40893919,  76975.30495554,  39697.9349351 ,
        69519.83095145,  76975.30495554,  69519.83095145,   9876.03891874,
        39697.9349351 ,  54608.88294328,  62064.35694737,  69519.83095145,
        84430.77895963,  62064.35694737,  76975.30495554,  84430.77895963,
        84430.77895963,  69519.83095145,  62064.35694737, -12490.38309353,
        76975.30495554,  62064.35694737,  69519.83095145,  62064.35694737,
        47153.40893919,  84430.77895963,  69519.83095145,  47153.40893919,
        62064.35694737,  69519.83095145,  69519.83095145,  76975.30495554,
        84430.77895963,  47153.40893919,  47153.40893919,  91886.25296372,
       -19945.85709762,  69519.83095145,  91886.25296372,  62064.35694737,
        62064.35694737,  39697.9349351 ,  47153.40893919,  76975.30495554,
        62064.35694737,  62064.35694737,  69519.83095145,  62064.35694737,
```

```

62064.35694737, 76975.30495554, 62064.35694737, 69519.83095145,
47153.40893919, 62064.35694737, 76975.30495554, 84430.77895963,
39697.9349351 , 47153.40893919, 91886.25296372, 84430.77895963,
84430.77895963, 69519.83095145, 69519.83095145, 84430.77895963,
69519.83095145, 76975.30495554, 47153.40893919, 32242.46093101,
62064.35694737, 17331.51292283, 76975.30495554, 39697.9349351 ,
24786.98692692, 39697.9349351 , 54608.88294328, 76975.30495554,
54608.88294328, 84430.77895963, 84430.77895963, 54608.88294328,
39697.9349351 , 69519.83095145, 39697.9349351 , 17331.51292283,
62064.35694737, 76975.30495554, 54608.88294328, 39697.9349351 ,
76975.30495554, 76975.30495554, 84430.77895963, 24786.98692692,
54608.88294328, 69519.83095145, 47153.40893919, 62064.35694737,
62064.35694737, 62064.35694737, 54608.88294328, 69519.83095145,
91886.25296372, 39697.9349351 , 69519.83095145, 24786.98692692,
24786.98692692, 91886.25296372, 69519.83095145, 54608.88294328,
62064.35694737, 84430.77895963, 84430.77895963, 54608.88294328,
76975.30495554, 54608.88294328, 54608.88294328, 9876.03891874,
84430.77895963, 76975.30495554, 47153.40893919, 32242.46093101,
69519.83095145, 17331.51292283, 76975.30495554, 76975.30495554,
24786.98692692, 91886.25296372, 62064.35694737, 84430.77895963,
69519.83095145, 62064.35694737, 84430.77895963, 69519.83095145,
39697.9349351 , 84430.77895963, 76975.30495554, 69519.83095145,
47153.40893919, 24786.98692692, 32242.46093101, 47153.40893919,
84430.77895963, 84430.77895963, 62064.35694737, 47153.40893919,
47153.40893919, 76975.30495554, 69519.83095145, 76975.30495554,
62064.35694737, 2420.56491465, 39697.9349351 , 39697.9349351 ,
84430.77895963, 32242.46093101, 62064.35694737, 62064.35694737,
84430.77895963, 54608.88294328, 54608.88294328, 32242.46093101,
39697.9349351 , 84430.77895963, 69519.83095145, -5034.90908944,
32242.46093101, 84430.77895963, 47153.40893919, 2420.56491465,
76975.30495554, 62064.35694737, 47153.40893919, 54608.88294328,
69519.83095145, 32242.46093101, 62064.35694737, 76975.30495554,
91886.25296372, 91886.25296372, 54608.88294328, 54608.88294328,
84430.77895963, 76975.30495554, 32242.46093101, 54608.88294328,
91886.25296372, 62064.35694737, 62064.35694737, 47153.40893919,
62064.35694737, 62064.35694737, 76975.30495554, 54608.88294328,
54608.88294328, 84430.77895963, 47153.40893919, 15633.57798773,
91886.25296372, 69519.83095145, 24786.98692692, 62064.35694737,
24786.98692692, 62064.35694737, 24786.98692692, 91886.25296372,
76975.30495554, 39697.9349351 , 76975.30495554, 84430.77895963,
69519.83095145, 84430.77895963, 69519.83095145, 69519.83095145,
76975.30495554, 2420.56491465, 17331.51292283, 24786.98692692,
84430.77895963, 62064.35694737, 39697.9349351 , 39697.9349351 ,
91886.25296372, 54608.88294328, 62064.35694737, 54608.88294328,
69519.83095145, 69519.83095145, 84430.77895963, 54608.88294328,
54608.88294328, 69519.83095145, 76975.30495554, 54608.88294328.

```

y\_pred.shape

(251,)

```
from sklearn.metrics import mean_absolute_error, mean_absolute_percentage_error, r2_score, mean_
```

```
mean_absolute_error(y_test, y_pred)
```

28689.77223478409

```
r2_score(y_test,y_pred)
```

0.06731083195243104

```
mean_squared_error(y_test,y_pred)
```

3701406312.046656


```
mean_absolute_percentage_error(y_test,y_pred)
```

0.5629069619670382

```
mp.scatter(y_test,y_pred)
```

```
mp.title('Actual prices vs Predicted prices')
```

```
mp.show()
```



```
df1_new=df2.sample(1)
```

```
df1_new.shape
```

(1, 8)

financial market news

```
df3=pd.read_csv(r'https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/Financial%20M
```

```
df3.head()
```

	Date	Label	News 1	News 2	News 3	News 4	News 5	News 6	News 7
0	01-01-2010	0	McIlroy's men catch cold from Gudjonsson	Obituary: Brian Walsh	Workplace blues leave employers in the red	Classical review: Rattle	Dance review: Merce Cunningham	Genetic tests to be used in setting premiums	Opera review: Bohème
1	02-01-2010	0	Warning from history points to crash	Investors flee to dollar haven	Banks and tobacco in favour	Review: Llama Farmers	War jitters lead to sell-off	Your not-so-secret history	Review: The Northumbrian Sinfonia
2	03-01-2010	0	Comment: Why Israel's peaceniks feel betrayed	Court deals blow to seizure of drug assets	An ideal target for spooks	World steps between two sides intent on war	What the region's papers say	Comment: Fear and rage in Palestine	Povey's resentment fuels Palestinian fury
3	04-01-2010	1	£750,000-a-goal Weah aims parting shot	Newcastle pay for Fletcher years	Brown sent to the stands for Scotland qualifier	Tourists wary of breaking new ground	Canary Wharf climbs into the FTSE 100	Review: Bill Bailey	Review: Classical
4	05-01-2010	1	Leeds arrive in Turkey to the silence of the fans	One woman's vision offers loan lifeline	Working Lives: How world leaders worked	Working Lives: Tricks of the trade	Working Lives: six-hour days, long lunches and...	Pop review: We Love UK	Working Lives: review: Mar Mo

5 rows × 27 columns



```
df3.tail()
```

	Date	Label	News 1	News 2	News 3	News 4	News 5	News 6
4096	20-03-2021	0	Barclays and RBS shares suspended from trading...	Pope says Church should ask forgiveness from g...	Poland 'shocked' by xenophobic abuse of Poles ...	There will be no second referendum, cabinet ag...	Scotland welcome to join EU, Merkel ally says	Sterling dips below Friday's 31-year low amid ...
4097	21-03-2021	1	2,500 Scientists To Australia: If You Want To ...	The personal details of 112,000 French police ...	S&P cuts United Kingdom sovereign credit r...	Huge helium deposit found in Africa	CEO of the South African state broadcaster qui...	Brexit cost investors \$2 trillion, the worst o...
4098	22-03-2021	1	Explosion At Airport In Istanbul	Yemeni former president: Terrorism is the offs...	UK must accept freedom of movement to access E...	Devastated: scientists too late to captive bre...	British Labor Party leader Jeremy Corbyn loses...	A Muslim Shop in the UK Was Just Firebombed Wh...
4099	23-03-2021	1	Jamaica proposes marijuana dispensers for tour...	Stephen Hawking says pollution and 'stupidity'...	Boris Johnson says he will not run for Tory pa...	Six gay men in Ivory Coast were abused and for...	Switzerland denies citizenship to Muslim immig...	Palestinian terrorist stabs israeli teen girl ...
			A 117-	IMF chief	The			

```
df3.describe()
```

	Label
count	4101.000000
mean	0.528164
std	0.499267
min	0.000000
25%	0.000000
50%	1.000000
75%	1.000000
max	1.000000

```
df3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4101 entries, 0 to 4100
Data columns (total 27 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Date        4101 non-null   object
1   Label       4101 non-null   int64
2   News 1     4101 non-null   object
3   News 2     4101 non-null   object
4   News 3     4101 non-null   object
5   News 4     4101 non-null   object
6   News 5     4101 non-null   object
7   News 6     4101 non-null   object
8   News 7     4101 non-null   object
9   News 8     4101 non-null   object
10  News 9     4101 non-null   object
11  News 10    4101 non-null   object
12  News 11    4101 non-null   object
13  News 12    4101 non-null   object
14  News 13    4101 non-null   object
15  News 14    4101 non-null   object
16  News 15    4101 non-null   object
17  News 16    4101 non-null   object
18  News 17    4101 non-null   object
19  News 18    4101 non-null   object
20  News 19    4101 non-null   object
21  News 20    4101 non-null   object
22  News 21    4101 non-null   object
23  News 22    4101 non-null   object
24  News 23    4100 non-null   object
25  News 24    4098 non-null   object
26  News 25    4098 non-null   object
dtypes: int64(1), object(26)
memory usage: 865.2+ KB
```

```
df3.shape
```

```
(4101, 27)
```

```
df3.columns
```

```
Index(['Date', 'Label', 'News 1', 'News 2', 'News 3', 'News 4', 'News 5',
      'News 6', 'News 7', 'News 8', 'News 9', 'News 10', 'News 11', 'News 12',
      'News 13', 'News 14', 'News 15', 'News 16', 'News 17', 'News 18',
      'News 19', 'News 20', 'News 21', 'News 22', 'News 23', 'News 24',
      'News 25'],
      dtype='object')
```

```
df3.index
```



```
RangeIndex(start=0, stop=4101, step=1)
```

```
len(df3.index)
```

```
4101
```

```
news=[]
```

```
for row in range(0,len(df3.index)):
```

```
    news.append(' '.join(str(x) for x in df3.iloc[row,2:27]))
```

```
type(news)
```

```
list
```

```
news[0]
```

```
'McIlroy's men catch cold from Gudjonsson Obituary: Brian Walsh Workplace blues leave e  
mployers in the red Classical review: Rattle Dance review: Merce Cunningham Genetic tes  
ts to be used in setting premiums Opera review: La Bohème Pop review: Britney Spears Th  
eatre review: The Circle Wales face a fraught night Under-21 round-up Smith off to blo  
t his copybook Finns taking the mickey Praise wasted as Brown studies injury options Ir  
eland warw of minnows Finland 0 - 0 England Healy a marked man Hannu birthday Harners &
```

```
x=news
```

```
type(x)
```

```
list
```

```
from sklearn.feature_extraction.text import CountVectorizer
```

```
cv=CountVectorizer(lowercase=True,ngram_range=(1,1))
```

```
x=cv.fit_transform(x)
```

```
x.shape
```

```
(4101, 48527)
```

```
y=df3['Label']
```

```
y.shape
```

(4101,)

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.4,stratify=y,random_state=252
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf=RandomForestClassifier(n_estimators=200)
```

```
rf.fit(x_train,y_train)
```

```
RandomForestClassifier(n_estimators=200)
```

```
y_pred=rf.predict(x_test)
```

```
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
```

```
confusion_matrix(y_test,y_pred)
```

```
array([[310, 851],
       [316, 984]])
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.50	0.27	0.35	1161
1	0.54	0.76	0.63	1300
accuracy			0.53	2461
macro avg	0.52	0.51	0.49	2461
weighted avg	0.52	0.53	0.50	2461

```
df4=pd.read_csv(r"https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/Movies%20Reco
```

```
df4.head()
```

	Movie_ID	Movie_Title	Movie_Genre	Movie_Language	Movie_Budget	Movie_Popularity	
0	1	Four Rooms	Crime Comedy	en	4000000	22.876230	
1	2	Star Wars	Adventure Action Science Fiction	en	11000000	126.393695	
2	3	Finding Nemo	Animation Family	en	94000000	85.688789	
3	4	Forrest Gump	Comedy Drama Romance	en	55000000	138.133331	
4	5	American	Drama	en	150000000	80.878605	

df4.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4760 entries, 0 to 4759
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	Movie_ID	4760 non-null	int64
1	Movie_Title	4760 non-null	object
2	Movie_Genre	4760 non-null	object
3	Movie_Language	4760 non-null	object
4	Movie_Budget	4760 non-null	int64
5	Movie_Popularity	4760 non-null	float64
6	Movie_Release_Date	4760 non-null	object
7	Movie_Revenue	4760 non-null	int64
8	Movie_Runtime	4758 non-null	float64
9	Movie_Vote	4760 non-null	float64
10	Movie_Vote_Count	4760 non-null	int64
11	Movie_Homepage	1699 non-null	object
12	Movie_Keywords	4373 non-null	object
13	Movie_Overview	4757 non-null	object
14	Movie_Production_House	4760 non-null	object
15	Movie_Production_Country	4760 non-null	object

```
16  Movie_Spoken_Language    4760 non-null  object
17  Movie_Tagline            3942 non-null  object
18  Movie_Cast               4733 non-null  object
19  Movie_Crew               4760 non-null  object
20  Movie_Director           4738 non-null  object
dtypes: float64(3), int64(4), object(14)
memory usage: 781.1+ KB
```

df4.shape

```
(4760, 21)
```

df4.columns

```
Index(['Movie_ID', 'Movie_Title', 'Movie_Genre', 'Movie_Language',
      'Movie_Budget', 'Movie_Popularity', 'Movie_Release_Date',
      'Movie_Revenue', 'Movie_Runtime', 'Movie_Vote', 'Movie_Vote_Count',
      'Movie_Homepage', 'Movie_Keywords', 'Movie_Overview',
      'Movie_Production_House', 'Movie_Production_Country',
      'Movie_Spoken_Language', 'Movie_Tagline', 'Movie_Cast', 'Movie_Crew',
      'Movie_Director'],
      dtype='object')
```

```
df_features=df4[['Movie_ID', 'Movie_Title', 'Movie_Genre', 'Movie_Language']].fillna('')
```

df\_features.shape

```
(4760, 4)
```

df\_features

	Movie_ID	Movie_Title	Movie_Genre	Movie_Language
0	1	Four Rooms	Crime Comedy	en

```
x=df_features['Movie_Genre']+' '+df_features['Movie_Title']+' '+df_features['Movie_Language']
```

```
2 3 Finding Nemo Animation Family en
```

```
x
```

```
0 Crime Comedy Four Rooms en
1 Adventure Action Science Fiction Star Wars en
2 Animation Family Finding Nemo en
3 Comedy Drama Romance Forrest Gump en
4 Drama American Beauty en
```

```
...
```

```
4755 Horror Midnight Cabaret en
4756 Comedy Family Drama Growing Up Smith en
4757 Thriller Drama 8 Days en
4758 Family Running Forever en
4759 Documentary To Be Frank, Sinatra at 100 en
```

```
Length: 4760, dtype: object
```

```
4760 rows x 4 columns
```

```
x.shape
```

```
(4760,)
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
tf=TfidfVectorizer()
```

```
x=tf.fit_transform(x)
```

```
x.shape
```

```
(4760, 4678)
```

```
print(x)
```

```
(0, 1367) 0.09007714202335212
(0, 3486) 0.7429127511854295
(0, 1613) 0.5912191932404761
(0, 869) 0.1708587827138408
(0, 976) 0.24745483330221305
(1, 4487) 0.6508001682997612
(1, 3902) 0.550439231907864
(1, 1509) 0.2814439744740556
(1, 3582) 0.2817769790975338
(1, 86) 0.21313740258007258
(1, 102) 0.2462864718546055
(1, 1367) 0.09371359348263858
```

```

(2, 2862)    0.6732410589493643
(2, 1523)    0.6200613847496372
(2, 1467)    0.24661415234213488
(2, 208)     0.30787730152862025
(2, 1367)    0.08162954584662571
(3, 1850)    0.6742499796285102
(3, 1602)    0.6742499796285102
(3, 3475)    0.2057784569495271
(3, 1238)    0.13304095730802998
(3, 1367)    0.08175187608144899
(3, 869)     0.15506737578583457
(4, 401)     0.7754819374908652
(4, 177)     0.5986761264571135
:           :
(4755, 2027) 0.2706809764473197
(4755, 1367) 0.08948840682655666
(4756, 1833) 0.6172310324083766
(4756, 4359) 0.4518371988005941
(4756, 3790) 0.5684756204008815
(4756, 1238) 0.121790151891555
(4756, 1467) 0.2260971844085065
(4756, 1367) 0.07483840771174535
(4756, 869)  0.14195387369812426
(4757, 1072) 0.8961993131347343
(4757, 4150) 0.3345007373814785
(4757, 1238) 0.2483049922847849
(4757, 1367) 0.15258007286186046
(4758, 3518) 0.6750368818728584
(4758, 1597) 0.6750368818728584
(4758, 1467) 0.28266019908870965
(4758, 1367) 0.09356082553006834
(4759, 2)    0.4622024356803555
(4759, 3737) 0.4622024356803555
(4759, 1622) 0.44084567731400226
(4759, 389)  0.36782643816611105
(4759, 1185) 0.25065051674667876
(4759, 289)  0.335899903341829
(4759, 4187) 0.2535775056782416
(4759, 1367) 0.05604140509889691

```

```
from sklearn.metrics.pairwise import cosine_similarity
```

```
ss=cosine_similarity(x)
```

```
ss
```

```

array([[1.          , 0.00844145, 0.00735296, ..., 0.01374398, 0.00842769,
        0.00504805],
       [0.00844145, 1.          , 0.0076498 , ..., 0.01429883, 0.00876792,
        0.00525184],
       [0.00735296, 0.0076498 , 1.          , ..., 0.01245504, 0.07734533,
        0.00457463],

```

```
...,
[0.01374398, 0.01429883, 0.01245504, ..., 1.          , 0.01427552,
 0.0085508 ],
[0.00842769, 0.00876792, 0.07734533, ..., 0.01427552, 1.          ,
 0.00524328],
[0.00504805, 0.00525184, 0.00457463, ..., 0.0085508 , 0.00524328,
 1.          ]])
```

```
ss.shape
```

```
(4760, 4760)
```

```
favourite_movie_name=input('enter movie name:')
```

```
enter movie name:avatar
```

```
movie_title_list=df4['Movie_Title'].tolist()
```

```
import difflib
```

```
movie_recommendation=difflib.get_close_matches(favourite_movie_name,movie_title_list)
print(movie_recommendation)
```

```
['Avatar']
```

```
index_of_close_matchmovie=df4['Movie_ID'].values[0]
print(index_of_close_matchmovie)
```

```
1
```

```
recommended_score=list(enumerate(ss[index_of_close_matchmovie]))
print(recommended_score)
```

```
[(0, 0.008441452669654322), (1, 1.0000000000000002), (2, 0.00764979807564309), (3, 0.007
```



```
len(recommended_score)
```

```
4760
```

```
sorted_similar_movies=sorted(recommended_score,key=lambda x:x[1],reverse=True)
print(sorted_similar_movies)
```

```
[(1, 1.0000000000000002), (4676, 0.7418973411852932), (601, 0.6207065009413768), (2233,
```



```
print('top 10 suggested movies : \n')
```

```

i=1
for movie in sorted_similar_movies:
    index=movie[0]
    title_from_index=df4[df4.index==index]['Movie_Title'].values[0]
    if(i<=10):
        print(i,')', title_from_index)
        i+=1

```

top 10 suggested movies :

- 1 ) Star Wars
- 2 ) Star Wars: Clone Wars: Volume 1
- 3 ) Star Wars: Episode I - The Phantom Menace
- 4 ) Star Trek
- 5 ) Star Wars: Episode III - Revenge of the Sith
- 6 ) Star Wars: Episode II - Attack of the Clones
- 7 ) U.F.O.
- 8 ) Star Trek Beyond
- 9 ) Star Trek: Generations
- 10 ) Star Trek: Insurrection

```
df5=pd.read_csv("https://github.com/YBI-foundation/Dataset/raw/main/WhiteWineQuality.csv",sep
```

```
df5.head()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulp
<b>0</b>	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	
<b>1</b>	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	
<b>2</b>	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	
<b>3</b>	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	

```
df5.tail()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	s
<b>4893</b>	6.2	0.21	0.29	1.6	0.039	24.0	92.0	0.99114	3.27	
<b>4894</b>	6.6	0.32	0.36	8.0	0.047	57.0	168.0	0.99490	3.15	
<b>4895</b>	6.5	0.24	0.19	1.2	0.041	30.0	111.0	0.99254	2.99	
<b>4896</b>	5.5	0.29	0.30	1.1	0.022	20.0	110.0	0.98869	3.34	



```
df5.describe()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	density
<b>count</b>	4898.000000	4898.000000	4898.000000	4898.000000	4898.000000	4898.000000	4898.000000
<b>mean</b>	6.854788	0.278241	0.334192	6.391415	0.045772	35.308085	138.000000
<b>std</b>	0.843868	0.100795	0.121020	5.072058	0.021848	17.007137	42.000000
<b>min</b>	3.800000	0.080000	0.000000	0.600000	0.009000	2.000000	9.000000
<b>25%</b>	6.300000	0.210000	0.270000	1.700000	0.036000	23.000000	108.000000
<b>50%</b>	6.800000	0.260000	0.320000	5.200000	0.043000	34.000000	134.000000
<b>75%</b>	7.300000	0.320000	0.390000	9.900000	0.050000	46.000000	167.000000

```
df5.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4898 entries, 0 to 4897
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          4898 non-null   float64
1   volatile acidity       4898 non-null   float64
2   citric acid            4898 non-null   float64
3   residual sugar         4898 non-null   float64
4   chlorides              4898 non-null   float64
5   free sulfur dioxide    4898 non-null   float64
6   total sulfur dioxide   4898 non-null   float64
7   density                4898 non-null   float64
8   pH                    4898 non-null   float64
9   sulphates              4898 non-null   float64
10  alcohol                4898 non-null   float64
11  quality                4898 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 459.3 KB
```

```
df5.columns
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
      'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
      'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')
```

```
df5.shape
```

```
(4898, 12)
```

```
df5['quality'].value_counts()
```

```
6    2198
5    1457
7     880
8     175
4     163
3      20
9       5
Name: quality, dtype: int64
```

```
df5.groupby('quality').mean()
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density
quality								
3	7.600000	0.333250	0.336000	6.392500	0.054300	53.325000	170.600000	0.99488
4	7.129448	0.381227	0.304233	4.628221	0.050098	23.358896	125.279141	0.99421
5	6.933974	0.302011	0.337653	7.334969	0.051546	36.432052	150.904598	0.99526
6	6.837671	0.260564	0.338025	6.441606	0.045217	35.650591	137.047316	0.99396
7	6.734716	0.262767	0.325625	5.186477	0.038191	34.125568	125.114773	0.99241
8	6.657143	0.277400	0.326514	5.671429	0.038314	36.720000	126.165714	0.99221

```
y=df5['quality']
```

```
y.shape
```

```
(4898,)
```

```
y
```

```
0    6
1    6
2    6
3    6
4    6
..
4893  6
4894  5
4895  6
4896  7
4897  6
Name: quality, Length: 4898, dtype: int64
```

```
x=df5[['density','pH','sulphates','alcohol']]
```

```
x=df5.drop(['quality'], axis=1)
```

```
x.shape
```

```
(4898, 11)
```

```
x
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	s
<b>0</b>	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.00100	3.00	
<b>1</b>	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.99400	3.30	
<b>2</b>	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.99510	3.26	
<b>3</b>	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	
<b>4</b>	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	
...	...	...	...	...	...	...	...	...	...	...
<b>4893</b>	6.2	0.21	0.29	1.6	0.039	24.0	92.0	0.99114	3.27	
<b>4894</b>	6.6	0.32	0.36	8.0	0.047	57.0	168.0	0.99490	3.15	
<b>4895</b>	6.5	0.24	0.19	1.2	0.041	30.0	111.0	0.99254	2.99	
<b>4896</b>	5.5	0.29	0.30	1.1	0.022	20.0	110.0	0.98869	3.34	
<b>4897</b>	6.0	0.21	0.38	0.8	0.020	22.0	98.0	0.98941	3.26	

```
from sklearn.preprocessing import StandardScaler
```

```
ss=StandardScaler()
```

```
x=ss.fit_transform(x)
```

```
x
```

```
array([[ 1.72096961e-01, -8.17699008e-02,  2.13280202e-01, ...,
        -1.24692128e+00, -3.49184257e-01, -1.39315246e+00],
       [-6.57501128e-01,  2.15895632e-01,  4.80011213e-02, ...,
        7.40028640e-01,  1.34184656e-03, -8.24275678e-01],
       [ 1.47575110e+00,  1.74519434e-02,  5.43838363e-01, ...,
```

```

4.75101984e-01, -4.36815783e-01, -3.36667007e-01],
...,
[-4.20473102e-01, -3.79435433e-01, -1.19159198e+00, ...,
-1.31315295e+00, -2.61552731e-01, -9.05543789e-01],
[-1.60561323e+00, 1.16673788e-01, -2.82557040e-01, ...,
1.00495530e+00, -9.62604939e-01, 1.85757201e+00],
[-1.01304317e+00, -6.77100966e-01, 3.78559282e-01, ...,
4.75101984e-01, -1.48839409e+00, 1.04489089e+00]])

from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.4,random_state=2529)

x_train.shape,x_test.shape,y_train.shape,y_test.shape

((1959, 11), (2939, 11), (1959,), (2939,))

from sklearn.svm import SVC

svc=SVC()

svc.fit(x_train,y_train)

SVC()

y_pred=svc.predict(x_test)

y_pred.shape

(2939,)

y_pred

array([6, 5, 6, ..., 6, 6, 5])

from sklearn.metrics import confusion_matrix,classification_report

print(confusion_matrix(y_test,y_pred))

[[ 0  0  6  6  0  0  0]
 [ 0  1 62 34  4  0  0]
 [ 0  0 499 350  4  0  0]
 [ 0  0 273 982 61  0  0]
 [ 0  0  23 389 132  0  0]
 [ 0  0  1  83  28  0  0]
 [ 0  0  0  0  1  0  0]]

```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	12
4	1.00	0.01	0.02	101
5	0.58	0.58	0.58	853
6	0.53	0.75	0.62	1316
7	0.57	0.24	0.34	544
8	0.00	0.00	0.00	112
9	0.00	0.00	0.00	1
accuracy			0.55	2939
macro avg	0.38	0.23	0.22	2939
weighted avg	0.55	0.55	0.51	2939

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1318: Undefined
_warn_prf(average, modifier, msg_start, len(result))
```

```
y=df5['quality'].apply(lambda y_value: 1 if y_value>=10 else 0)
```

```
y.value_counts()
```

```
0    4898
Name: quality, dtype: int64
```

```
df2_new=df5.sample(1)
```

```
df2_new
```

fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	s
------------------	---------------------	----------------	-------------------	-----------	---------------------------	----------------------------	---------	----	---

```
df2_new.shape
```

```
(1, 12)
```

```
x_new=df2_new.drop(['quality'],axis=1)
```

```
x_new=ss.fit_transform(x_new)
```

```
y_pred_new=svc.predict(x_new)
```

```
y_pred_new
```

```
array([6])
```

big sales prediction

```
df6=pd.read_csv(r"https://raw.githubusercontent.com/YBI-Foundation/Dataset/main/Big%20Sales%20Dataset.csv")
```

```
df6.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP
0	FDT36	12.3	Low Fat	0.111448	Baking Goods	33.4874
1	FDT36	12.3	Low Fat	0.111904	Baking Goods	33.9874
2	FDT36	12.3	LF	0.111728	Baking Goods	33.9874
3	FDT36	12.3	Low Fat	0.000000	Baking Goods	34.3874
4	FDP12	9.8	Regular	0.045523	Baking Goods	35.0874

```
df6.tail()
```

Item\_Identifier Item\_Weight Item\_Fat\_Content Item\_Visibility Item\_Type Item

Starchv

df6.describe()

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
<b>count</b>	11815.000000	14204.000000	14204.000000	14204.000000	14204.000000
<b>mean</b>	12.788355	0.065953	141.004977	1997.830681	218.000000
<b>std</b>	4.654126	0.051459	62.086938	8.371664	182.000000
<b>min</b>	4.555000	0.000000	31.290000	1985.000000	3.000000
<b>25%</b>	8.710000	0.027036	94.012000	1987.000000	92.000000
<b>50%</b>	12.500000	0.054021	142.247000	1999.000000	176.000000
<b>75%</b>	16.750000	0.094037	185.855600	2004.000000	298.000000
<b>max</b>	30.000000	0.328391	266.888400	2009.000000	3122.000000

df6.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Item_Identifier                       14204 non-null  object
1   Item_Weight                           11815 non-null  float64
2   Item_Fat_Content                       14204 non-null  object
3   Item_Visibility                       14204 non-null  float64
4   Item_Type                             14204 non-null  object
5   Item_MRP                              14204 non-null  float64
6   Outlet_Identifier                     14204 non-null  object
7   Outlet_Establishment_Year             14204 non-null  int64
8   Outlet_Size                           14204 non-null  object
9   Outlet_Location_Type                  14204 non-null  object
10  Outlet_Type                           14204 non-null  object
11  Item_Outlet_Sales                     14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

df6.shape

(14204, 12)

df6.columns

```
Index(['Item_Identifier', 'Item_Weight', 'Item_Fat_Content', 'Item_Visibility',
      'Item_Type', 'Item_MRP', 'Outlet_Identifier',
      'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
      'Outlet_Type', 'Item_Outlet_Sales'],
      dtype='object', name='columns')
```

```
'Outlet_Establishment_Year', 'Outlet_Size', 'Outlet_Location_Type',
'Outlet_Type', 'Item_Outlet_Sales'],
dtype='object')
```

```
df6['Item_Weight'].fillna(df6.groupby(['Item_Type'])['Item_Weight'].transform('mean'), inplace
```

```
df6.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14204 entries, 0 to 14203
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Item_Identifier        14204 non-null  object
1   Item_Weight            14204 non-null  float64
2   Item_Fat_Content       14204 non-null  object
3   Item_Visibility        14204 non-null  float64
4   Item_Type              14204 non-null  object
5   Item_MRP               14204 non-null  float64
6   Outlet_Identifier      14204 non-null  object
7   Outlet_Establishment_Year 14204 non-null  int64
8   Outlet_Size            14204 non-null  object
9   Outlet_Location_Type   14204 non-null  object
10  Outlet_Type            14204 non-null  object
11  Item_Outlet_Sales      14204 non-null  float64
dtypes: float64(4), int64(1), object(7)
memory usage: 1.3+ MB
```

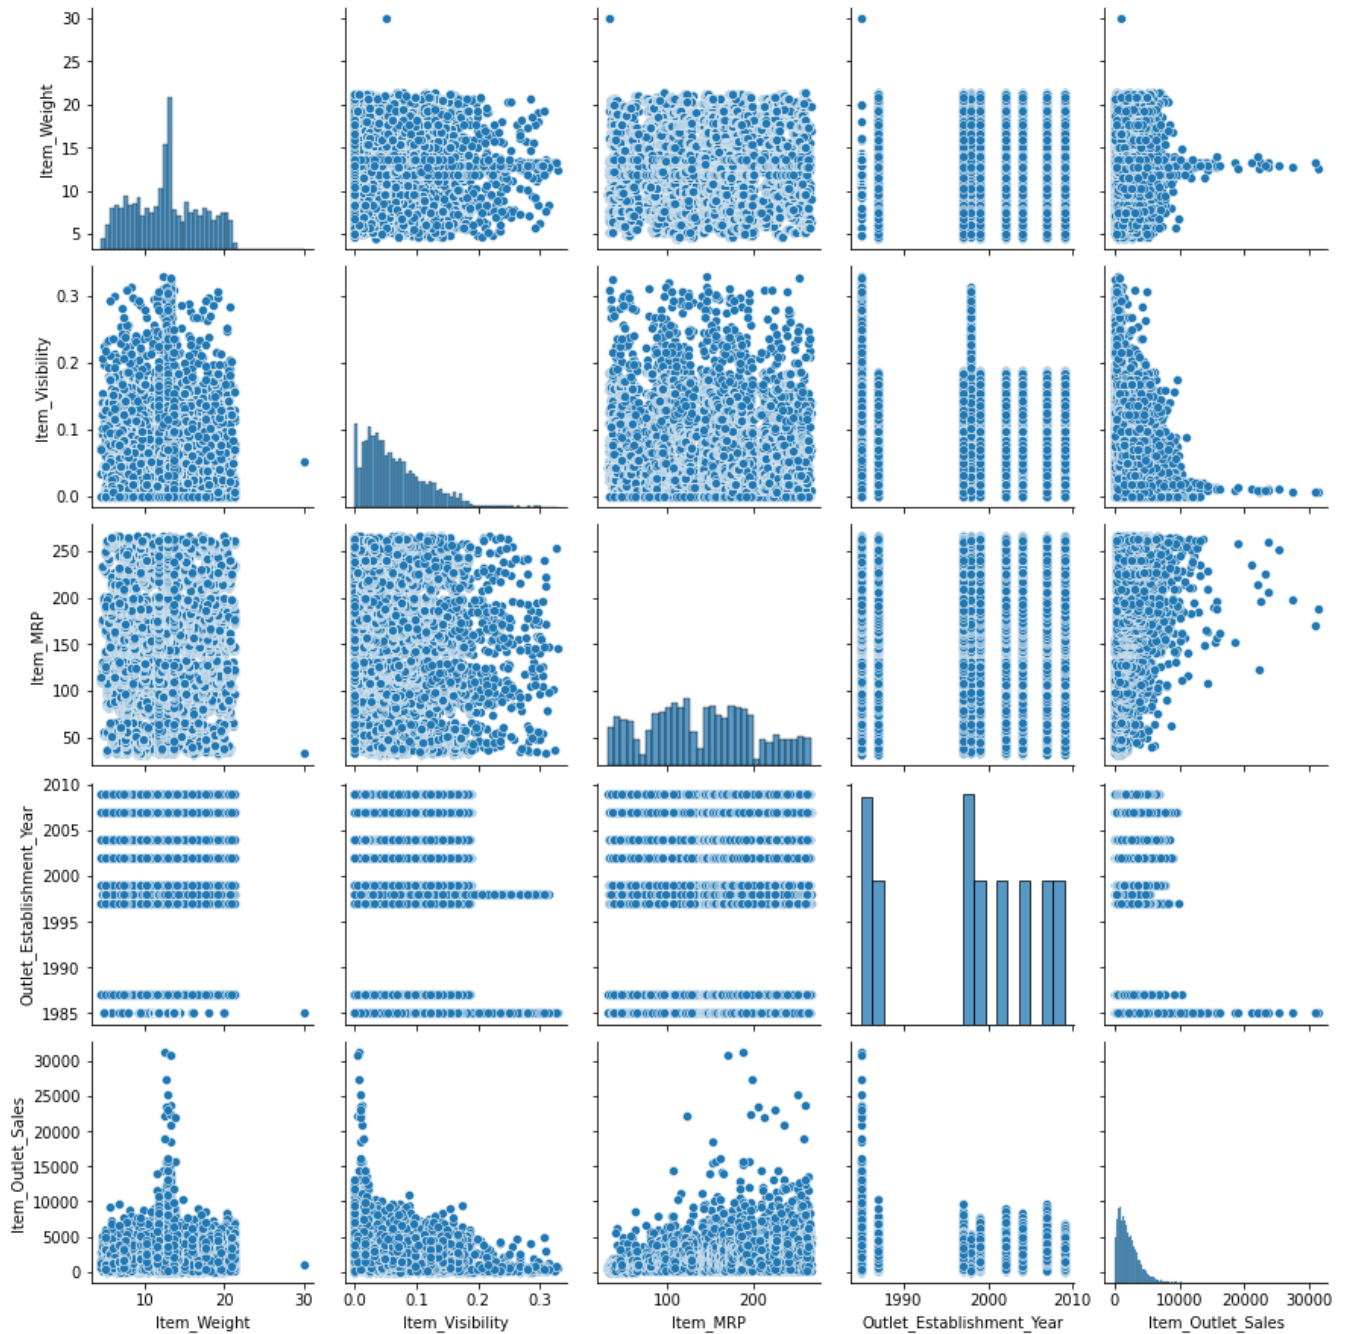
```
df6.describe()
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	14204.000000	14204.000000	14204.000000	14204.000000	14204.000000
mean	12.790642	0.065953	141.004977	1997.830681	216.343846
std	4.251186	0.051459	62.086938	8.371664	186.834946
min	4.555000	0.000000	31.290000	1985.000000	16.129000
25%	9.300000	0.027036	94.012000	1987.000000	91.269000
50%	12.800000	0.054021	142.247000	1999.000000	176.669000
75%	16.000000	0.094037	185.855600	2004.000000	291.269000
max	30.000000	0.328391	266.888400	2009.000000	3126.270000

```
sn.pairplot(df6)
```



<seaborn.axisgrid.PairGrid at 0x7fe585354cd0>



```
df6['Item_Weight'].value_counts()
```

```
13.194406    346
12.867687    335
13.332012    261
```

```

12.566713    250
13.238358    195
...
5.860000      7
7.850000      6
9.035000      6
4.615000      6
30.000000     1
Name: Item_Weight, Length: 432, dtype: int64

```

```
df6['Item_Fat_Content'].value_counts()
```

```

Low Fat    8485
Regular    4824
LF         522
reg        195
low fat    178
Name: Item_Fat_Content, dtype: int64

```

```
df6.replace({'Item_Fat_Content':{'Low Fat':0,'Regular':1}},inplace=True)
```

```
df6['Item_Fat_Content'].value_counts()
```

```

0         8485
1         4824
LF         522
reg        195
low fat    178
Name: Item_Fat_Content, dtype: int64

```

```
df6['Item_Type'].value_counts()
```

```

Fruits and Vegetables    2013
Snack Foods              1989
Household                1548
Frozen Foods             1426
Dairy                   1136
Baking Goods            1086
Canned                  1084
Health and Hygiene       858
Meat                    736
Soft Drinks             726
Breads                  416
Hard Drinks             362
Others                  280
Starchy Foods           269
Breakfast               186
Seafood                 89
Name: Item_Type, dtype: int64

```

```
df6.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP
0	FDT36	12.3	0	0.111448	Baking Goods	33.4874
1	FDT36	12.3	0	0.111904	Baking Goods	33.9874
2	FDT36	12.3	LF	0.111728	Baking Goods	33.9874
3	FDT36	12.3	0	0.000000	Baking Goods	34.3874
4	FDP12	9.8	1	0.045523	Baking Goods	35.0874

```
y=df6['Item_Outlet_Sales']
```

```
y
```

```
0      436.608721
1      443.127721
2      564.598400
3     1719.370000
4      352.874000
...
14199   4984.178800
14200   2885.577200
14201   2885.577200
14202   3803.676434
14203   3644.354765
Name: Item_Outlet_Sales, Length: 14204, dtype: float64
```

```
y.shape
```

```
(14204,)
```

```
x=df6[['Item_Type','Item_Weight','Item_MRP',"Outlet_Size",'Outlet_Type']]
```

```
x=df6.drop(['Item_Identifier'],axis=1)
```

```
x.shape
```

```
(14204, 11)
```

```
x
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Iden
<b>0</b>	12.300000	0	0.111448	Baking Goods	33.4874	C
<b>1</b>	12.300000	0	0.111904	Baking Goods	33.9874	C
<b>2</b>	12.300000	LF	0.111728	Baking Goods	33.9874	C
<b>3</b>	12.300000	0	0.000000	Baking Goods	34.3874	C
<b>4</b>	9.800000	1	0.045523	Baking Goods	35.0874	C
...	...	...	...	...	...	...
<b>14199</b>	12.800000	0	0.069606	Starchy Foods	261.9252	C
<b>14200</b>	12.800000	0	0.070013	Starchy Foods	262.8252	C
<b>14201</b>	12.800000	0	0.069561	Starchy Foods	263.0252	C
<b>14202</b>	13.659758	0	0.069282	Starchy Foods	263.5252	C
<b>14203</b>	12.800000	0	0.069727	Starchy Foods	263.6252	C

14204 rows × 11 columns



```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
```

```
x_std=df6[['Item_Weight','Item_MRP','Item_Visibility']]
```

```
x_std=sc.fit_transform(x_std)
```

```
x_std
```

```
array([[ -0.11541705, -1.73178716,  0.88413635],
       [ -0.11541705, -1.72373366,  0.89300616],
       [ -0.11541705, -1.72373366,  0.88958331],
```


```
...,
[ 0.00220132, 1.96538148, 0.07011952],
[ 0.20444792, 1.97343499, 0.06469366],
[ 0.00220132, 1.97504569, 0.07334891]])
```

```
x[['Item_Weight', 'Item_MRP', 'Item_Visibility']] = pd.DataFrame(x_std, columns=['Item_Weight', 'I
```

```
x
```

	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Iden
<b>0</b>	-0.115417	0	0.884136	Baking Goods	-1.731787	(
<b>1</b>	-0.115417	0	0.893006	Baking Goods	-1.723734	(
<b>2</b>	-0.115417	LF	0.889583	Baking Goods	-1.723734	(
<b>3</b>	-0.115417	0	-1.281712	Baking Goods	-1.717291	(
<b>4</b>	-0.703509	1	-0.397031	Baking Goods	-1.706016	(
...	...	...	...	...	...	
<b>14199</b>	0.002201	0	0.070990	Starchy Foods	1.947664	(
<b>14200</b>	0.002201	0	0.078898	Starchy Foods	1.962160	(
<b>14201</b>	0.002201	0	0.070120	Starchy Foods	1.965381	(
<b>14202</b>	0.204448	0	0.064694	Starchy Foods	1.973435	(
<b>14203</b>	0.002201	0	0.073349	Starchy Foods	1.975046	(

14204 rows × 11 columns



```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.2,random_state=2529)
```

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
((2840, 11), (11364, 11), (2840,), (11364,))
```

```
from sklearn.linear_model import LogisticRegression

lr=LogisticRegression()

lr.fit(x_train,y_train)

LogisticRegression()

y_pred=lr.predict(x_test)

from sklearn.svm import SVC

sc=SVC()

from sklearn.metrics import mean_squared_error,mean_absolute_error,r2_score

mean_absolute_error(y_test,y_pred)

0.42438130155820347


mean_squared_error(y_test,y_pred)

0.42438130155820347

r2_score(y_test,y_pred)

-0.6975266323890179

mp.scatter(y_test,y_pred)
mp.title('Actual prices vs Predicted prices')
mp.show()
```




hill and valley prediction

```
df7=pd.read_csv("https://github.com/YBI-foundation/Dataset/raw/main/Hill%20Valley%20Dataset.c
```

```
df7.head()
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9
0	39.02	36.49	38.20	38.85	39.38	39.74	37.02	39.53	38.81
1	1.83	1.71	1.77	1.77	1.68	1.78	1.80	1.70	1.75
2	68177.69	66138.42	72981.88	74304.33	67549.66	69367.34	69169.41	73268.61	74465.84
3	44889.06	39191.86	40728.46	38576.36	45876.06	47034.00	46611.43	37668.32	40980.89
4	5.70	5.40	5.28	5.38	5.27	5.61	6.00	5.38	5.34


5 rows × 101 columns



```
df7.tail()
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9
1207	13.00	12.87	13.27	13.04	13.19	12.53	14.31	13.33	13.63
1208	48.66	50.11	48.55	50.43	50.09	49.67	48.95	48.65	48.63
1209	10160.65	9048.63	8994.94	9514.39	9814.74	10195.24	10031.47	10202.28	9152.99
1210	34.81	35.07	34.98	32.37	34.16	34.03	33.31	32.48	35.63
1211	8489.43	7672.98	9132.14	7985.73	8226.85	8554.28	8838.87	8967.24	8635.14

5 rows × 101 columns



```
df7.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1212 entries, 0 to 1211
Columns: 101 entries, V1 to Class
dtypes: float64(100), int64(1)
memory usage: 956.5 KB
```

```
df7.describe()
```

	V1	V2	V3	V4	V5	
count	1212.000000	1212.000000	1212.000000	1212.000000	1212.000000	1212.0
mean	8169.091881	8144.306262	8192.653738	8176.868738	8128.297211	8173.0
std	17974.950461	17881.049734	18087.938901	17991.903982	17846.757963	17927.1
min	0.920000	0.900000	0.850000	0.890000	0.880000	0.8
25%	19.602500	19.595000	18.925000	19.277500	19.210000	19.5
50%	301.425000	295.205000	297.260000	299.720000	295.115000	294.3
75%	5358.795000	5417.847500	5393.367500	5388.482500	5321.987500	5328.0
max	117807.870000	108896.480000	119031.350000	110212.590000	113000.470000	116848.3

8 rows × 101 columns

◀ ▶

```
print(df7.columns.tolist())
```

```
['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14']
```

◀ [REDACTED] ▶

```
df7.shape
```

 $(1212, 101)$ 

```
df7['Class'].value_counts()
```

```
0      606
1      606
Name: Class, dtype: int64
```

```
df7.groupby('Class').mean()
```

	V1	V2	V3	V4	V5	V6	
Class							
0	7913.333251	7825.339967	7902.497294	7857.032079	7775.610198	7875.436337	7804.
1	8424.850512	8463.272558	8482.810182	8496.705396	8480.984224	8470.623680	8572.

2 rows × 100 columns

◀ ▶

```
y=df7[ 'Class' ]
```



```
y.shape
```

```
(1212,)
```

```
y
```

```
0      0
1      1
2      1
3      0
4      0
..
1207    1
1208    0
1209    1
1210    1
1211    0
```

```
Name: Class, Length: 1212, dtype: int64
```

```
x=df7[['V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', "V8", "V9", "V10"]]
```

```
x=df7.drop('Class',axis=1)
```

```
x.shape
```


```
(1212, 100)
```

```
x
```

	V1	V2	V3	V4	V5	V6	V7	V8	
0	39.02	36.49	38.20	38.85	39.38	39.74	37.02	39.53	38.85
1	1.83	1.71	1.77	1.77	1.68	1.78	1.80	1.70	1.77


```
mp.plot(x.iloc[0,:])
mp.title('valley')
```

```
Text(0.5, 1.0, 'valley')
```



```
mp.plot(x.iloc[1,:])
mp.title('Hill')
```

```
Text(0.5, 1.0, 'Hill')
```



```
from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
```

```
x=sc.fit_transform(x)
```

x

```
array([[ -0.45248681, -0.45361784, -0.45100881, ..., -0.45609618,
        -0.45164274, -0.45545496],
       [ -0.45455665, -0.45556372, -0.45302369, ..., -0.45821768,
        -0.45362255, -0.45755405],
       [  3.33983504,  3.24466709,  3.58338069, ...,  3.5427869 ,
        3.27907378,  3.74616847],
       ...,
       [  0.11084204,  0.0505953 ,  0.04437307, ...,  0.12533312,
        0.04456025,  0.06450317],
       [ -0.45272112, -0.45369729, -0.45118691, ..., -0.45648861,
        -0.45190136, -0.45569511],
       [  0.01782872, -0.02636986,  0.05196137, ...,  0.03036056,
        0.01087365,  0.03123129]])
```

x.shape

```
(1212, 100)
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=0.1,stratify=y,random_state=253
```

```
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
((121, 100), (1091, 100), (121,), (1091,))
```

```
from sklearn.linear_model import LogisticRegression
```

```
lr=LogisticRegression()
```

```
lr.fit(x_train,y_train )
```

```
LogisticRegression()
```

```
y_pred=lr.predict(x_test)
```

```
y_pred.shape
```

```
(1091,)
```

```
y_pred
```

```
array([1, 1, 1, ..., 1, 1, 1])
```

```
lr.predict_proba(x_test)
```

```
array([[0.49339312, 0.50660688],
       [0.49226548, 0.50773452],
       [0.49340597, 0.50659403],
       ...,
       [0.49339979, 0.50660021],
       [0.49342196, 0.50657804],
       [0.49339765, 0.50660235]])
```

```
from sklearn.metrics import confusion_matrix,classification_report
```

```
print(confusion_matrix(y_test,y_pred))
```

```
[[101 445]
 [ 18 527]]
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.85	0.18	0.30	546
1	0.54	0.97	0.69	545
accuracy			0.58	1091
macro avg	0.70	0.58	0.50	1091
weighted avg	0.70	0.58	0.50	1091

```
x1_new=df7.sample(2)
```

```
x1_new
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...
<b>1202</b>	1.86	1.93	1.91	1.90	1.97	1.89	2.02	1.87	2.02	2.03	...
<b>612</b>	409.53	393.75	388.72	388.45	413.10	406.14	415.73	413.29	411.99	411.11	... 41

2 rows × 101 columns

```
x1_new.shape
```

```
(2, 101)
```

```
x1_new=x1_new.drop('Class',axis=1)
```

x1\_new

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...
<b>1202</b>	1.86	1.93	1.91	1.90	1.97	1.89	2.02	1.87	2.02	2.03	...
<b>612</b>	409.53	393.75	388.72	388.45	413.10	406.14	415.73	413.29	411.99	411.11	... 38

2 rows × 100 columns

x1\_new

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	...
<b>1202</b>	1.86	1.93	1.91	1.90	1.97	1.89	2.02	1.87	2.02	2.03	...
<b>612</b>	409.53	393.75	388.72	388.45	413.10	406.14	415.73	413.29	411.99	411.11	... 38

2 rows × 100 columns

```
x1_new=ss.fit_transform(x1_new)
```

```
y_pred_new=lr.predict(x1_new)
```

y\_pred\_new

```
array([1, 1, 1, ..., 1, 1, 1])
```

```
lr.predict_proba(x1_new)
```

```
array([[0.49299481, 0.50700519],
       [0.49193574, 0.50806426],
       [0.49300773, 0.50699227],
       ...,
       [0.49300159, 0.50699841],
       [0.49302497, 0.50697503],
       [0.49299931, 0.50700069]])
```